
DS551/CS551/CS525 2024 Fall

Project 3 - Deep Q-learning

— 10/6/2024 —

Outline

- Introduction
 - Game Playing : Breakout
- Deep Reinforcement Learning
 - Deep Q-Learning (DQN)
 - Improvements to DQN
- Grading & Format
 - Grading Policy
 - Code Format
 - Submission
- WPI Turing or Google Cloud Platform & Pytorch Tutorial

Introduction

Environment

Breakout




- Get average reward ≥ 40 in 100 episodes (5 lives per episode)
- In testing, we consider each episode with its all 5 lives
- With **OpenAI's Atari wrapper** (modified by us a little bit)

Deep Reinforcement Learning


Deep Q-Learning (DQN)

“classic” deep Q-learning algorithm:

- 
1. take some action \mathbf{a}_i and observe $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$, add it to \mathcal{B} Replay buffer
 2. sample mini-batch $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j\}$ from \mathcal{B} uniformly
 3. compute $y_j = r_j + \gamma \max_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mathbf{a}'_j)$ using *target* network $Q_{\phi'}$
 4. $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_\phi(\mathbf{s}_j, \mathbf{a}_j) - y_j)$ Fixed target-Q
 5. update ϕ' : copy ϕ every N steps

Introduction

Training Plot

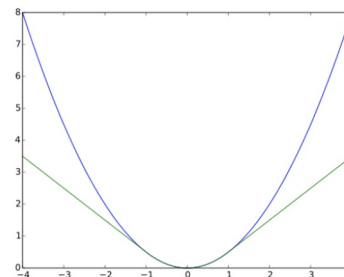


- X-axis : number of training steps
- Y-axis : average reward in every 30 **lives** (not 30 complete episodes).

Deep Reinforcement Learning

Deep Q-Learning (DQN)

- The action should act ϵ -greedily
 - Random action with probability ϵ
- Linearly decay ϵ from 1.0 to some small value, say 0.025
 - Decay per step: $(\epsilon - \epsilon_{\min}) / \text{number of epsilon step}$
- Hyperparameters (just suggestion)
 - Replay Buffer Memory Size 10,000 (*deque*)
 - Start to train DQN with buffer size 5000
 - Update Target Network every 5000 steps
 - Learning Rate $1.5e-4$, Batch Size 32
 - Adam optimizer
 - Huber Loss (*F.smooth_l1_loss*)
 - Clip gradients between (-1,1)



Green is the Huber loss and blue is the quadratic loss (Wikipedia)

$$L_{\delta}(a) = \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta, \\ \delta(|a| - \frac{1}{2}\delta), & \text{otherwise.} \end{cases}$$

Deep Reinforcement Learning

Improvements to DQN

- Double Q-Learning
- Dueling Network
- Prioritized Replay Memory
- Noisy DQN
- Distributional DQN

<https://arxiv.org/pdf/1710.02298.pdf>

Deep Reinforcement Learning

Other Training Tips

- [How to use Pytorch](#)
- [Official DQN Pytorch Tutorial](#)
- [DQN Tutorial on Medium](#)
- [Official DQN paper](#)
- [See more tips on project website](#)
- <https://github.com/Illyyyt123/WPI-DS551-Fall23/blob/main/Project3/README.md>

Grading & Format

Grading Policy

- Python code (20 points)
- Trained Model (50 points)
 - Get averaging reward ≥ 40 in 100 episodes (each of 5 lives) in **Breakout**
 - With **OpenAI's Atari wrapper**
- PDF Report (30 points)
 - Describe your DQN model
 - Screenshot of the average score in 100 episodes
 - Plot the training curve (*training steps can be defined by yourself*)
 - X-axis: number of training steps
 - Y-axis: average reward in last 30 lives

Grading & Format

Code Format

- Please download all the .py files from project [github page](#)
- Follow the instructions in README to install packages
- **Six** functions you should implement in [agent_dqn.py](#)
 1. `__init__(self, env, args)`
 2. `init_game_setting(self)`
 3. `make_action(self, state, test)`
 4. `train(self)`
 5. `push(self)`
 6. `repaly_buffer(self)`
- **DO NOT** add any parameter in `__init__()`, `init_game_setting()` and `make_action()`
- You can change the seed
- You can add new functions in the [agent_dqn.py](#)

Grading & Format

Code Format

- **Two** functions you should implement in `dqn_model.py`
 1. `__init__(self)`
 2. `forward(self, x)`
- You can add parameters in these two functions
- You can add new functions in the `dqn_model.py`
- You can add your arguments in `argument.py` (if needed)
- Please do not change `test.py`, `main.py`, `environment.py`, `atari_wrapper.py` and `agent.py`

Grading & Format

Deliverables

- Deadline: **Tuesday Oct 29, 2024 (23:59)**
- Your submission **MUST** have following files
 - `agent_dqn.py`, `dqn_model.py`, `argument.py`
 - `[saved_model_file]` (.pth file)
 - `report.pdf`
 - README (with details of what files you have modified.)
 - other files you need
- If your model is too large for canvas, upload it to a cloud space (like dropbox, google drive) and provide the link to download the model

Grading & Format

Package

- Please use Python3
- The TA will execute 'python main.py --test_dqn' to run your code on **ubuntu+GPU**
- The execution for the model should be done within 20 minutes, excluding model download
- Allowed packages
 - a. PyTorch
 - b. Numpy
 - c. Scipy
 - d. Pandas
 - e. Python Standard Lib
 - f. etc.

Setup

- Recommended programming IDE (integrated development environment): VS code (See [install VS code](#))
- Install [Miniconda](#)
- Install [Python 3](#), by default, it's Python 3.11.4.
- For more details, please refer to project 3 website

<https://github.com/UrbanIntelligence/WPI-DS551-Fall24/tree/main/Project3>

Environment Preparation

- GPU resources:
 1. How to use WPI Turing GPUs with your WPI account
 2. Google Cloud <https://cloud.google.com/gpu>

backup

